

Lab 8: Machine Learning 1

Taylor Darby

10/22/2021

Clustering methods

Kmeans clustering in R is done with the “kmeans()” function. Here we makeup some data to test and learn with.

```
# 'rnorm' provides a random dataset distributed normally, and specifying x and y sets the number of rows
tmp <- c(rnorm(30, 3), rnorm(30, -3))

# 'cbind' binds the data in the columns together
data <- cbind(x=tmp, y=rev(tmp))

# We now have a dataset to perform k-means with
data
```


```
##           x           y
## [1,]  2.7550259 -4.1195588
## [2,]  3.0633328 -2.4829510
## [3,]  3.1881915 -4.3939522
## [4,]  1.1127519 -3.2394942
## [5,]  3.2874358 -4.7007221
## [6,]  1.6708457 -1.4849593
## [7,]  2.3694659 -0.4289242
## [8,]  2.0173767 -2.1843387
## [9,]  2.5870592 -2.4793419
## [10,] 1.8960701 -2.1617600
## [11,] 2.2677314 -2.9885094
## [12,] 2.8324203 -3.5932365
## [13,] 3.2959694 -2.0079145
## [14,] 3.8705091 -2.2604250
## [15,] 3.1937913 -2.9653960
## [16,] 2.9117002 -3.6942136
## [17,] 2.7912476 -2.9917906
## [18,] 3.2955767 -4.5235896
## [19,] 2.4762455 -4.1766312
## [20,] 2.6987888 -2.7648218
## [21,] 3.1586320 -1.8469826
## [22,] 2.2219785 -4.3122930
## [23,] 1.0754638 -2.3513653
## [24,] 1.4805525 -3.4124065
## [25,] 3.1304136 -2.2143707
```

##	[26,]	2.4905842	-3.1523270
##	[27,]	2.3265662	-2.8597214
##	[28,]	4.2469035	-2.6387041
##	[29,]	4.7027752	-3.5195426
##	[30,]	3.0637949	-1.7861700
##	[31,]	-1.7861700	3.0637949
##	[32,]	-3.5195426	4.7027752
##	[33,]	-2.6387041	4.2469035
##	[34,]	-2.8597214	2.3265662
##	[35,]	-3.1523270	2.4905842
##	[36,]	-2.2143707	3.1304136
##	[37,]	-3.4124065	1.4805525
##	[38,]	-2.3513653	1.0754638
##	[39,]	-4.3122930	2.2219785
##	[40,]	-1.8469826	3.1586320
##	[41,]	-2.7648218	2.6987888
##	[42,]	-4.1766312	2.4762455
##	[43,]	-4.5235896	3.2955767
##	[44,]	-2.9917906	2.7912476
##	[45,]	-3.6942136	2.9117002
##	[46,]	-2.9653960	3.1937913
##	[47,]	-2.2604250	3.8705091
##	[48,]	-2.0079145	3.2959694
##	[49,]	-3.5932365	2.8324203
##	[50,]	-2.9885094	2.2677314
##	[51,]	-2.1617600	1.8960701
##	[52,]	-2.4793419	2.5870592
##	[53,]	-2.1843387	2.0173767
##	[54,]	-0.4289242	2.3694659
##	[55,]	-1.4849593	1.6708457
##	[56,]	-4.7007221	3.2874358
##	[57,]	-3.2394942	1.1127519
##	[58,]	-4.3939522	3.1881915
##	[59,]	-2.4829510	3.0633328
##	[60,]	-4.1195588	2.7550259

Run `kmeans()`: Set “`k`”(“centers”) to 2 nstart to 20. The thing with `kmeans` is that you have to tell it how many clusters you want (in this case = 2)

```
km <- kmeans(data, centers = 2, nstart = 20)
km
```

```
## K-means clustering with 2 clusters of sizes 30, 30  
##  
## Cluster means:  
##      x           y  
## 1 -2.924547  2.715973  
## 2  2.715973 -2.924547  
##  
## Clustering vector:  
## [1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1 1 1  
## [39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1  
##  
## Within cluster sum of squares by cluster:
```

Hierarchical Clustering

We will use the 'hclust()' function on the same data as before and see how this method works.


```
hc <- hclust(dist(data))
hc
```

```
##
## Call:
## hclust(d = dist(data))
##
## Cluster method   : complete
## Distance         : euclidean
## Number of objects: 60
```

'hclust' has a plot method

```
plot(hc, cex=0.5)
abline(h=7, col="#05DCBB")
```

Cluster Dendrogram



```
dist(data)
hclust (*, "complete")
```

To find our membership vector we need to “cut” the tree and for this we use the ‘`cutree()`’ function and tell it the height to cut at.


```
cutree(hc, h=7)
```

[illegible]

We can also use ‘cutree()’ and state the number of k clusters we want...

```
grps <- cutree(hc, k=2)
```

```
plot(data, col=grps)
```



Principal Component Analysis (PCA)

PCA is useful for visualizing key variance in datasets with high dimensionality.

PCA of UK food data

Import UK food dataset

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url)
head(x)
```

```
##           X England Wales Scotland N.Ireland
## 1      Cheese      105   103      103        66
## 2 Carcass_meat     245   227      242       267
## 3   Other_meat     685   803      750       586
## 4         Fish     147   160      122        93
## 5 Fats_and_oils     193   235      184       209
## 6       Sugars     156   175      147       139
```

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

```
## Complete the following code to find out how many rows and columns are in x?
dim(x)
```

```
## [1] 17 5
```

```
## Preview the first 6 rows
head(x)
```

```
##           X England Wales Scotland N.Ireland
## 1      Cheese      105   103      103       66
## 2 Carcass_meat     245   227      242      267
## 3   Other_meat     685   803      750      586
## 4        Fish     147   160      122       93
## 5 Fats_and_oils     193   235      184      209
## 6        Sugars     156   175      147      139
```

Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

```
# Note how the minus indexing works
# rownames(x) <- x[,1] "This is not a good way to code because you will lose a column every time you run"
# x <- x[,-1]

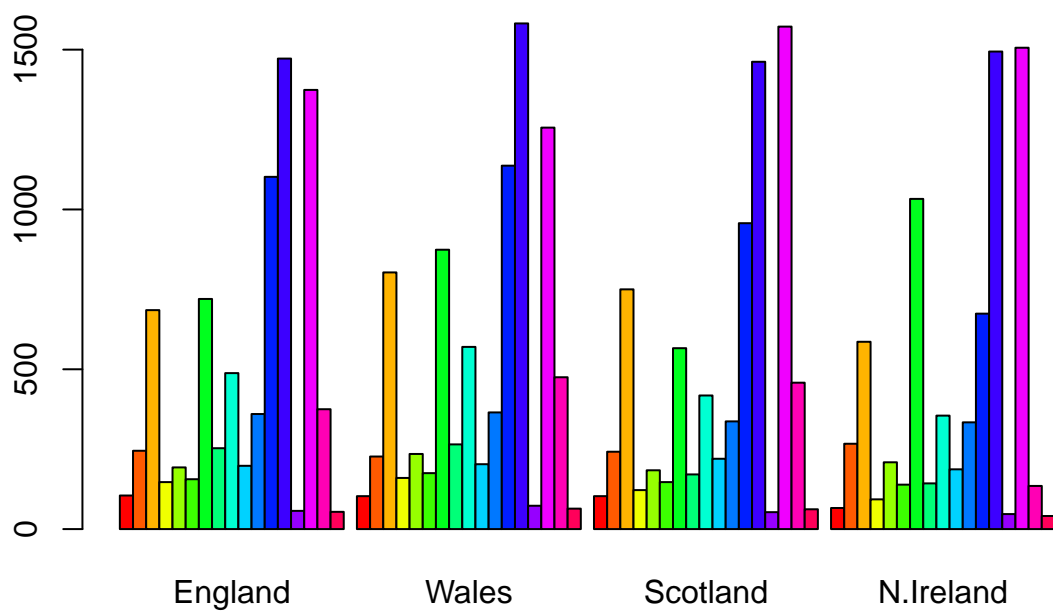
# This way is better
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names=1)
head(x)
```

```
##           England Wales Scotland N.Ireland
## Cheese      105   103      103       66
## Carcass_meat 245   227      242      267
## Other_meat   685   803      750      586
## Fish        147   160      122       93
## Fats_and_oils 193   235      184      209
## Sugars      156   175      147      139
```

Spotting major differences and trends


Let's plot the data

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



Q3: Changing what optional argument in the above `barplot()` function results in the following plot?


```
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```

Q5: Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

Make a pairwise plot

```
mycols <- rainbow( nrow(x))
pairs(x, col=mycols, pch=16)
```



Answer: The axis comparisons change across rows and columns. If the point lies on the diagonal then that food is consumed at the same rate in both countries.

Q6. What is the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

N. Ireland consumes less of one of the values than the rest of the countries.

PCA to the rescue


Here we will use the base R function for PCA, which is called 'prcomp()'. Note: This function wants the transpose of our data.

```
# Use the prcomp() PCA function
pca <- prcomp( t(x) )
summary(pca)
```

```
## Importance of components:
##              PC1      PC2      PC3      PC4
## Standard deviation 324.1502 212.7478 73.87622 4.189e-14
## Proportion of Variance 0.6744 0.2905 0.03503 0.000e+00
## Cumulative Proportion 0.6744 0.9650 1.00000 1.000e+00
```

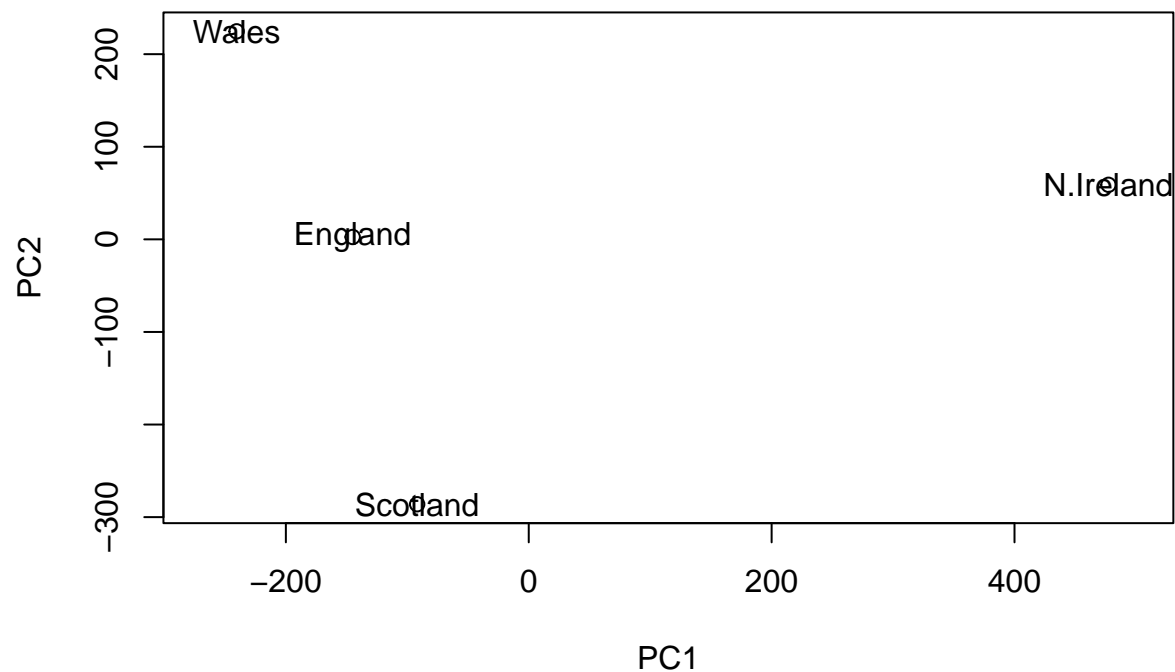
Plot this pca

```
plot(pca)
```




Q7. Complete the code below to generate a plot of PC1 vs PC2. The second line adds text labels over the data points.

```
# Plot PC1 vs PC2
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", xlim=c(-270,500))
text(pca$x[,1], pca$x[,2], colnames(x))
```



Q8. Customize your plot so that the colors of the country names match the colors in our UK and Ireland map and table at start of this document.

```
# Plot PC1 vs PC2
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", xlim=c(-270,500))
text(pca$x[,1], pca$x[,2], labels = colnames(x), col=c("orange", "red", "blue", "green"))
```



Below we can use the square of `pca$sdev`, which stands for “standard deviation”, to calculate how much variation in the original data each PC accounts for.

```
v <- round( pca$sdev^2/sum(pca$sdev^2) * 100 )
v
```

```
## [1] 67 29 4 0
```


```
## or the second row here...
```

```
z <- summary(pca)
z$importance
```

```
##              PC1      PC2      PC3      PC4
## Standard deviation 324.15019 212.74780 73.87622 4.188568e-14
## Proportion of Variance 0.67444 0.29052 0.03503 0.000000e+00
## Cumulative Proportion 0.67444 0.96497 1.00000 1.000000e+00
```


This information can be summarized in a plot of the variances (eigenvalues) with respect to the principal component number (eigenvector number), which is given below.

```
barplot(v, xlab="Principal Component", ylab="Percent Variation")
```



Digging deeper (variable loadings)

```
## Lets focus on PC1 as it accounts for > 90% of variance  
par(mar=c(10, 3, 0.35, 0))  
barplot( pca$rotation[,1], las=2 )
```



Q9: Generate a similar ‘loadings plot’ for PC2. What two food groups feature prominently and what does PC2 mainly tell us about?

One more PCA for today

Import RNAseq data

```
url2 <- "https://tinyurl.com/expression-CSV"
rna.data <- read.csv(url2, row.names=1)
head(rna.data)
```

```
##      wt1 wt2 wt3 wt4 wt5 ko1 ko2 ko3 ko4 ko5
## gene1 439 458 408 429 420 90 88 86 90 93
## gene2 219 200 204 210 187 427 423 434 433 426
## gene3 1006 989 1030 1017 973 252 237 238 226 210
## gene4 783 792 829 856 760 849 856 835 885 894
## gene5 181 249 204 244 225 277 305 272 270 279
## gene6 460 502 491 491 493 612 594 577 618 638
```


Q10: How many genes and samples are in this data set?

```
nrow(rna.data)
```

```
## [1] 100
```

```
## Again we have to take the transpose of our data
pca.rna <- prcomp(t(rna.data), scale=TRUE)
```

```
## Simple un polished plot of pc1 and pc2
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2")
```




```
summary(pca.rna)
```

```
## Importance of components:
##              PC1      PC2      PC3      PC4      PC5      PC6      PC7
## Standard deviation  9.6237  1.5198  1.05787  1.05203  0.88062  0.82545  0.80111
## Proportion of Variance 0.9262  0.0231  0.01119  0.01107  0.00775  0.00681  0.00642
## Cumulative Proportion 0.9262  0.9493  0.96045  0.97152  0.97928  0.98609  0.99251
##              PC8      PC9      PC10
## Standard deviation  0.62065  0.60342  3.348e-15
## Proportion of Variance 0.00385  0.00364  0.000e+00
## Cumulative Proportion 0.99636  1.00000  1.000e+00
```

Scree plot of RNA data

```
plot(pca.rna$x[,1:2])
text(pca.rna$x[,1:2], labels=colnames(rna.data))
```

OR Let's make the above scree plot ourselves and in so doing explore the object returned from `prcomp()` a little further. We can use the square of `pca$sdev`, which stands for “standard deviation”, to calculate how much variation in the original data each PC accounts for:

```
## Variance captured per PC
```

```
pca.var <- pca.rna$sdev^2
```

```
## Percent variance is often more informative to look at
```


```
pca.var.per <- round(pca.var/sum(pca.var)*100, 1)
```

```
pca.var.per
```

```
## [1] 92.6 2.3 1.1 1.1 0.8 0.7 0.6 0.4 0.4 0.0
```

We can use this to generate our own scree-plot like this

```
barplot(pca.var.per, main="Scree Plot",
        names.arg = paste0("PC", 1:10),
        xlab="Principal Component", ylab="Percent Variation")
```




Now lets make our main PCA plot a bit more attractive and useful

```
## A vector of colors for wt and ko samples
colvec <- colnames(rna.data)
colvec[grep("wt", colvec)] <- "red"
colvec[grep("ko", colvec)] <- "blue"

plot(pca.rna$x[,1], pca.rna$x[,2], col=colvec, pch=16,
     xlab=paste0("PC1 (", pca.var.per[1], "%)"),
     ylab=paste0("PC2 (", pca.var.per[2], "%)"))

text(pca.rna$x[,1], pca.rna$x[,2], labels = colnames(rna.data), pos=c(rep(4,5), rep(2,5)))
```



... left off on ggplot section